



Università degli Studi di Bari
Dipartimento di Informatica



LACAM
Machine Learning

Simplifying, Regularizing and Strengthening Sum-Product Network Structure Learning

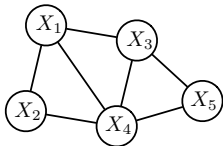
Antonio Vergari, Nicola Di Mauro and Floriana Esposito

August 22, 2015

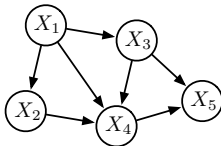
Summary

- ⊕ Sum-Product Networks refresher
- ⊕ Why and How Structure learning
- ⊕ Simplifying by limiting splits
- ⊕ Regulizing by effective early stopping
- ⊕ Strengthening by model averaging
- ⊕ Conclusions and further works

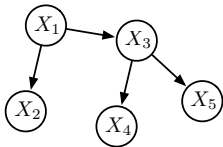
PGMs and Tractability



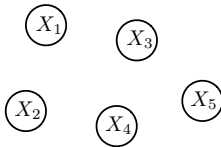
$$P(\mathbf{X}) = \frac{1}{Z} \prod_c \phi_c(\mathbf{X}_c)$$



$$P(\mathbf{X}) = \prod_{i=1}^n P(X_i | \mathbf{Pa}_i)$$

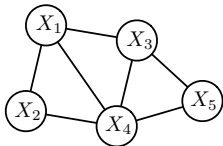


$$P(\mathbf{X}) = \prod_{i=1}^n P(X_i | Pa_i)$$

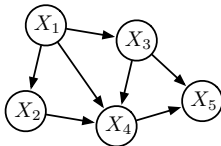


$$P(\mathbf{X}) = \prod_{i=1}^n P(X_i)$$

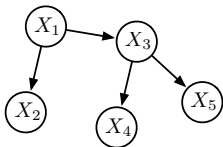
PGMs and Tractability



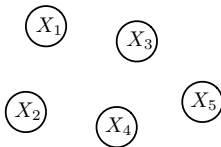
untractable



untractable





tractable



tractable

Sum-Product Networks (I)



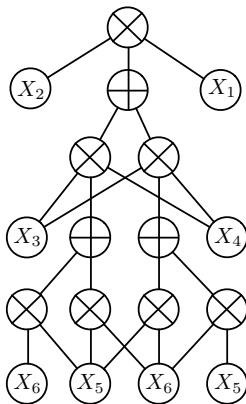
Compiling the partition function of a pdf into a **deep** architecture of **sum** and **product** nodes.

Product nodes define factorizations over independent vars, sum nodes mixtures. Leaves are tractable univariate distributions.

Products over nodes with different scopes (*decomposability*) and sums over nodes with same scopes (*completeness*) guarantee modeling a pdf (*validity*).

Considering only valid SPNs of *alternated layers of sum and products*.

Sum-Product Networks (II)



Bottom-up evaluation of the network:

$$S_{X_i}(x_j) = P(X_i = x_j)$$

$$S_+(\mathbf{x}) = \sum_{i \in ch(+)} w_i S_i(\mathbf{x})$$

$$S_\times(\mathbf{x}) = \prod_{i \in ch(\times)} S_i(\mathbf{x})$$

Inferences linear in the size of the network (# edges):

$$\oplus Z = S(*)$$

$$\oplus P(\mathbf{e}) = S(\mathbf{e})/S(*)$$

$$\oplus P(\mathbf{q}|\mathbf{e}) = \frac{P(\mathbf{q},\mathbf{e})}{P(\mathbf{e})} = \frac{S(\mathbf{q},\mathbf{e})}{S(\mathbf{e})}$$

$$\oplus MPE(\mathbf{q}, \mathbf{e}) = \max_{\mathbf{q}} P(\mathbf{q}, \mathbf{e}) = S^{max}(\mathbf{e})$$

How and Why Structure Learning

Fixed structures are hard to engineer and train (fully connected layers).

Automatic discovery of latent vars.

Constraint-based search formulation. Discover hidden variables for sum node mixtures and independences for product node components:

- ⊕ greedy top-down: KMeans on features [Dennis and Ventura 2012]; alternating clustering on instances and independence tests on features, **LearnSPN** [Gens and Domingos 2013]
- ⊕ greedy bottom up: merging feature regions by a *Bayesian-Dirichlet independence test*, and reducing edges by maximizing MI [Peharz, Geiger, and Pernkopf 2013]
- ⊕ **ID-SPN**: turning LearnSPN in log-likelihood guided expansion of sub-networks approximated by Arithmetic Circuits [Rooshenas and Lowd 2014]

Why Structure Quality Matters

Tractability is guaranteed if the network size is polynomial in # vars.

Comparing network sizes is better than comparing inference times.

Deeper networks are possibly more expressively efficient [Martens and Medabalimi 2014; Zhao, Melibari, and Poupart 2015]

Overcomplex networks do not generalize well

Structure quality desiderata: smaller but accurate, deeper but not wider, SPNs

LearnSPN (I)

Build a tree-like SPN by recursively split the data matrix:

- ⊕ splitting columns into pairs by a greedy **G Test** based procedure with threshold ρ :

$$G(X_i, X_j) = 2 \sum_{x_i \sim X_i} \sum_{x_j \sim X_j} c(x_i, x_j) \cdot \log \frac{c(x_i, x_j) \cdot |T|}{c(x_i)c(x_j)}$$

- ⊕ clustering instances with **online Hard-EM** with cluster penalty λ :

$$Pr(\mathbf{X}) = \sum_{C_i \in \mathbf{C}} \prod_{X_j \in \mathbf{X}} Pr(X_j | C_i) Pr(C_i)$$


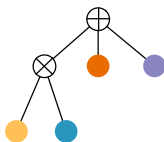
- ⊕ if there are less than m instances, put a **naive factorization** over leaves
- ⊕ each univariate distribution get **ML estimation** smoothed by α

	X_1	X_2	X_3	X_4	X_5
1					
2					
3					
4					
5					
6					
7					
8					

X_1	X_2	X_3	X_4	X_5
blue	blue	blue	blue	blue
blue	blue	blue	blue	blue
orange	orange	orange	orange	orange
orange	orange	orange	orange	orange
blue	blue	blue	blue	blue
purple	purple	purple	purple	purple
purple	purple	purple	purple	purple
purple	purple	purple	purple	purple

X_1	X_2	X_3	X_4	X_5
Orange	Orange	Orange	Blue	Blue
Orange	Orange	Orange	Blue	Blue
Orange	Orange	Orange	Orange	Orange
Orange	Orange	Orange	Orange	Orange
Orange	Orange	Orange	Blue	Blue
Blue	Blue	Blue	Blue	Blue
Blue	Blue	Blue	Blue	Blue
Blue	Blue	Blue	Blue	Blue

X_1	X_2	X_3	X_4	X_5
Orange	Orange	Orange	Blue	Blue
Orange	Orange	Orange	Blue	Blue
White	White	White	White	White
Orange	Orange	Orange	Blue	Blue
Purple	Purple	Purple	Purple	Purple
Purple	Purple	Purple	Purple	Purple



Simplifying by limiting node splits

LearSPN performs two interleaved *greedy hierarchical* divisive *clustering* processes (co-clustering).

Each process benefits from the other one improvements/highly suffers from other's mistakes.

Idea: slowing down the processes by limiting the number of nodes to split into. SPN-B, variant of LearnSPN that uses EM for mixture modeling with $k = 2$ to cluster rows.

Pros:

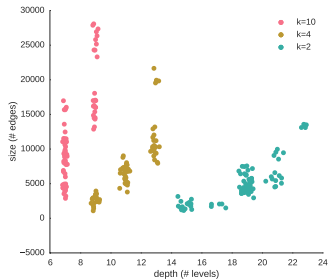
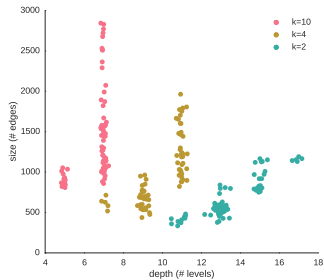
- ⊕ not committing to complex structures too early
- ⊕ same expressive power: successive splits allow for more node children
- ⊕ reducing node out fan increases the depth
- ⊕ same accuracy, smaller networks

Experimental setting

Classical setting for **generative** graphical models structure learning [Gens and Domingos 2013]:

- ⊕ comparing the **average log-likelihood** on predicting instances from a test set
- ⊕ 19 binary datasets from classification, recommendation, frequent pattern mining...[Lowd and Davis 2010] [Haaren and Davis 2012]
- ⊕ Training 75% Validation 10% Test 15% splits (no cv)
- ⊕ Model selection via *grid search* in the same parameter space:
 - ⊗ $\lambda \in \{0.2, 0.4, 0.6, 0.8\}$,
 - ⊗ $\rho \in \{5, 10, 15, 20\}$,
 - ⊗ $m \in \{1, 50, 100, 500\}$,
 - ⊗ $\alpha \in \{0.1, 0.2, 0.5, 1.0, 2.0\}$
- ⊕ comparing our variants against LearnSPN, ID-SPN and MT [Meilă and Jordan 2000]

Depth VS Size



Regularizing by effective early stopping

LearnSPN regularization is governed by α and m , however can be very ineffective:


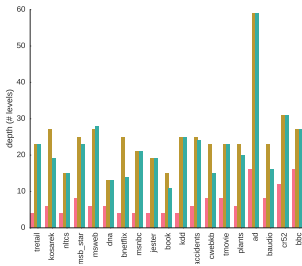
- ⊕ naive factorizations are too strong assumptions
- ⊕ best likelihood structures prefer smaller values for m to get accurate naive factorizations

Substituting naive factorizations with Bayesian trees as leaf distributions

$$P(\mathbf{X}) = \prod_j P(X_j | Pa_j):$$

- ⊕ learnable with Chow-Liu algorithm
- ⊕ still tractable multivariate distributions for marginals, conditionals and MPE
- ⊕ same or higher accuracy for larger values of m
- ⊕ possibly reducing structure complexity even more

Early stopping exp



Strengthening by model averaging

Interpreting sum nodes as *general additive estimators*. Leveraging classic statistical tools to learn them: **bagging**.

Draw k bootstrapped samples from the data, then grow an SPN S_{B_i} on each of them. Join them into a single SPN \hat{S} with a sum node:

$$\hat{S} = \sum_{i=1}^k \frac{1}{k} S_{B_i}$$

More robustness and less variance in the model.

Exponential number of nodes if done for each sum node (bootstrapping only at the root).

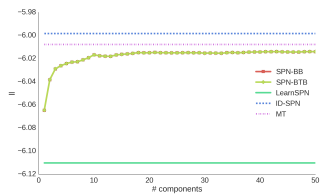
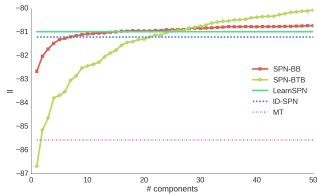
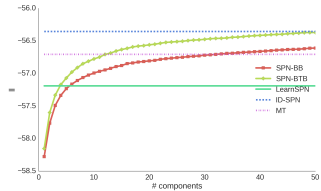
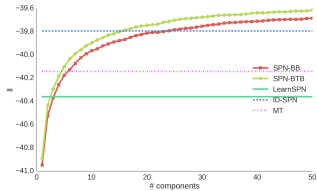
Two variants in the experiments: SPN-BB and SPN-BTB, whether Chow-Liu trees are employed or not.

LL exp

	LearnSPN	SPN-B	SPN-BT	ID-SPN	SPN-BB	SPN-BTB	MT
NLTCS	-6.110	-6.048	-6.048	-5.998	-6.014	-6.014	-6.008
MSNBC	-6.099	-6.040	-6.039	-6.040	-6.032	-6.033	-6.076
KDDCup2k	-2.185	-2.141	-2.141	-2.134	-2.122	-2.121	-2.135
Plants	-12.878	-12.813	-12.683	-12.537	-12.167	-12.089	-12.926
Audio	-40.360	-40.571	-40.484	-39.794	-39.685	-39.616	-40.142
Jester	-53.300	-53.537	-53.546	-52.858	-52.873	-53.600	-53.057
Netflix	-57.191	-57.730	-57.450	-56.355	-56.610	-56.371	-56.706
Accidents	-30.490	-29.342	-29.265	-26.982	-28.510	-28.351	-29.692
Retail	-11.029	-10.944	10.942	-10.846	-10.858	-10.858	-10.836
Pumsb-star	-24.743	-23.315	-23.077	-22.405	-22.866	-22.664	-23.702
DNA	-80.982	-81.913	-81.840	-81.211	-80.730	-80.068	-85.568
Kosarek	-10.894	-10.719	-10.685	-10.599	-10.690	-10.578	-10.615
MSWeb	-10.108	-9.833	-9.838	-9.726	-9.630	-9.614	-9.819
Book	-34.969	-34.306	-34.280	-34.136	-34.366	-33.818	-34.694
EachMovie	-52.615	-51.368	-51.388	-51.512	-50.263	-50.414	-54.513
WebKB	-158.164	-154.283	-153.911	-151.838	-151.341	-149.851	-157.001
Reuters-52	-85.414	-83.349	-83.361	-83.346	-81.544	-81.587	-86.531
BBC	-249.466	-247.301	-247.254	-248.929	-226.359	-226.560	-259.962
Ad	-19.760	-16.234	-15.885	-19.053	-13.785	-13.595	-16.012

Table : Average test log likelihoods for all algorithms.

Bagging exp



Conclusions and Further work

- ⊕ Structure quality evaluation matters
- ⊕ Deeper networks by applying a simplicity bias when splitting
- ⊕ Regularized SPNs by introducing Chow-Liu trees as leaves
- ⊕ More robust and accurate SPNs with bootstrapped sum nodes

References

- ⊕ Dennis, Aaron and Dan Ventura (2012). ``Learning the Architecture of Sum-Product Networks Using Clustering on Variables". In: *Advances in Neural Information Processing Systems 25*. Curran Associates, Inc., pp. 2033–2041 (cit. on p. 7).
- ⊕ Gens, Robert and Pedro Domingos (2013). ``Learning the Structure of Sum-Product Networks". In: *Proceedings of the 30th International Conference on Machine Learning*. JMLR Workshop and Conference Proceedings, pp. 873–880 (cit. on pp. 7, 12).
- ⊕ Haaren, Jan Van and Jesse Davis (2012). ``Markov Network Structure Learning: A Randomized Feature Generation Approach". In: *Proceedings of the 26th Conference on Artificial Intelligence*. AAAI Press (cit. on p. 12).
- ⊕ Lowd, Daniel and Jesse Davis (2010). ``Learning Markov Network Structure with Decision Trees". In: *Proceedings of the 10th IEEE International Conference on Data Mining*. IEEE Computer Society Press, pp. 334–343 (cit. on p. 12).
- ⊕ Martens, James and Venkatesh Medabalimi (2014). ``On the Expressive Efficiency of Sum Product Networks". In: *CoRR* abs/1411.7717 (cit. on p. 8).
- ⊕ Meilä, Marina and Michael I. Jordan (2000). ``Learning with mixtures of trees". In: *Journal of Machine Learning Research* 1, pp. 1–48 (cit. on p. 12).
- ⊕ Peharz, Robert, Bernhard Geiger, and Franz Pernkopf (2013). ``Greedy Part-Wise Learning of Sum-Product Networks". In: *Machine Learning and Knowledge Discovery in Databases*. Vol. 8189. LNCS. Springer, pp. 612–627 (cit. on p. 7).
- ⊕ Rooshenas, Amirmohammad and Daniel Lowd (2014). ``Learning Sum-Product Networks with Direct and Indirect Variable Interactions". In: *Proceedings of the 31st International Conference on Machine Learning*. JMLR Workshop and Conference Proceedings, pp. 710–718 (cit. on p. 7).
- ⊕ Zhao, Han, Mazen Melibari, and Pascal Poupart (2015). ``On the Relationship between Sum-Product Networks and Bayesian Networks". In: *CoRR* abs/1501.01239. URL: <http://arxiv.org/abs/1501.01239> (cit. on p. 8).