



Università degli Studi di Bari
Dipartimento di Informatica



LACAM
Machine Learning

Simplifying, Regularizing and Strengthening Sum-Product Network Structure Learning

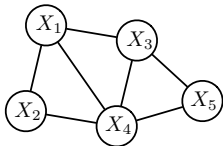
Antonio Vergari, Nicola Di Mauro and Floriana Esposito

August 27, 2015

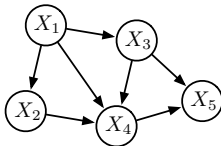
Summary

- ⊕ Sum-Product Networks refresher
- ⊕ Why and How Structure learning
- ⊕ Simplifying by limiting splits
- ⊕ Regulizing by effective early stopping
- ⊕ Strengthening by model averaging
- ⊕ Conclusions and further works

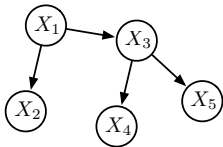
PGMs and Tractability



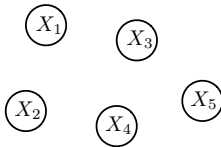
$$P(\mathbf{X}) = \frac{1}{Z} \prod_c \phi_c(\mathbf{X}_c)$$



$$P(\mathbf{X}) = \prod_{i=1}^n P(X_i | \mathbf{Pa}_i)$$

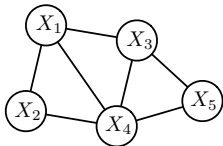


$$P(\mathbf{X}) = \prod_{i=1}^n P(X_i | Pa_i)$$

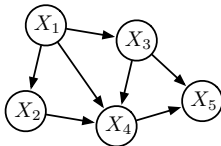


$$P(\mathbf{X}) = \prod_{i=1}^n P(X_i)$$

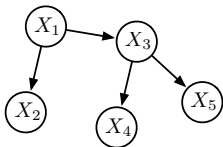
PGMs and Tractability



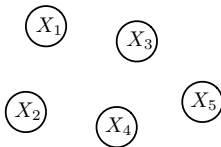
untractable



untractable



tractable

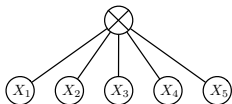


tractable

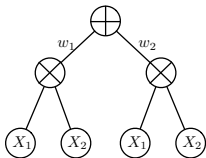
Sum-Product Networks (I)

SPNs are DAGs *compiling* the partition function of a joint pdf into a **deep** architecture of **sum** and **product** nodes.

Product nodes define factorizations over independent components, sum nodes weighted mixtures. Leaves are tractable univariate distributions.



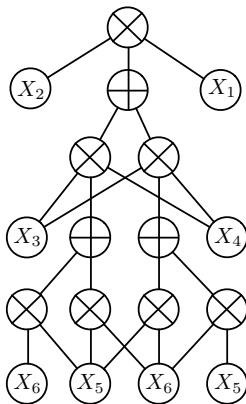
Products over nodes with different scopes (*decomposability*) and sums over nodes with same scopes (*completeness*) guarantee modeling a pdf (*validity*).



The size of the network is the number of edges in it.

The depth of the network is the number of alternated layers.

Sum-Product Networks (II)



Bottom-up evaluation of the network:

$$S_{X_i}(x_j) = P(X_i = x_j)$$

$$S_+(\mathbf{x}) = \sum_{i \in ch(+)} w_i S_i(\mathbf{x})$$

$$S_\times(\mathbf{x}) = \prod_{i \in ch(\times)} S_i(\mathbf{x})$$

Inferences linear in the size of the network (# edges):

$$\oplus Z = S(*)$$

$$\oplus P(\mathbf{e}) = S(\mathbf{e})/S(*)$$

$$\oplus P(\mathbf{q}|\mathbf{e}) = \frac{P(\mathbf{q},\mathbf{e})}{P(\mathbf{e})} = \frac{S(\mathbf{q},\mathbf{e})}{S(\mathbf{e})}$$

$$\oplus MPE(\mathbf{q}, \mathbf{e}) = \max_{\mathbf{q}} P(\mathbf{q}, \mathbf{e}) = S^{max}(\mathbf{e})$$

How and Why Structure Learning

Fixed structures are hard to engineer and train (fully connected layers).

Automatic discovery of latent vars.

Constraint-based search formulation. Discover hidden variables for sum node mixtures and independences for product node components:

- ⊕ greedy top-down: KMeans on features [Dennis and Ventura 2012]; alternating clustering on instances and independence tests on features, **LearnSPN** [Gens and Domingos 2013]
- ⊕ greedy bottom up: merging feature regions by a *Bayesian-Dirichlet independence test*, and reducing edges by maximizing MI [Peharz, Geiger, and Pernkopf 2013]
- ⊕ **ID-SPN**: turning LearnSPN in log-likelihood guided expansion of sub-networks approximated by Arithmetic Circuits [Rooshenas and Lowd 2014]

Why Structure Quality Matters

Tractability is guaranteed if the network size is polynomial in # vars.

Comparing network sizes is better than comparing inference times.

Deeper networks are possibly more expressively efficient [Martens and Medabalimi 2014; Zhao, Melibari, and Poupart 2015]

Overcomplex networks do not generalize well

Structure quality desiderata: smaller but accurate, deeper but not wider, SPNs

LearnSPN (I)

Build a tree-like SPN by recursively split the data matrix:

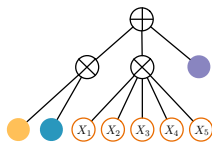
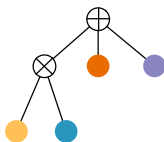
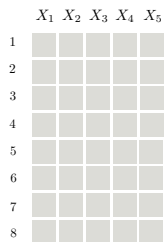
- ⊕ splitting columns into pairs by a greedy **G Test** based procedure with threshold ρ :

$$G(X_i, X_j) = 2 \sum_{x_i \sim X_i} \sum_{x_j \sim X_j} c(x_i, x_j) \cdot \log \frac{c(x_i, x_j) \cdot |T|}{c(x_i)c(x_j)}$$

- ⊕ clustering instances with **online Hard-EM** with cluster penalty λ :

$$Pr(\mathbf{X}) = \sum_{C_i \in \mathbf{C}} \prod_{X_j \in \mathbf{X}} Pr(X_j | C_i) Pr(C_i)$$

- ⊕ if there are less than m instances, put a **naive factorization** over leaves
- ⊕ each univariate distribution get **ML estimation** smoothed by α



Simplifying by limiting node splits

LearSPN performs two interleaved *greedy hierarchical* divisive *clustering* processes (co-clustering).

Each process benefits from the other one improvements/highly suffers from other's mistakes.

Idea: slowing down the processes by limiting the number of nodes to split into. SPN-B, variant of LearnSPN that uses EM for mixture modeling with $k = 2$ to cluster rows.

Pros:

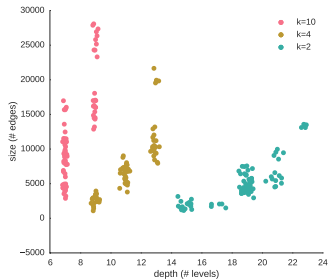
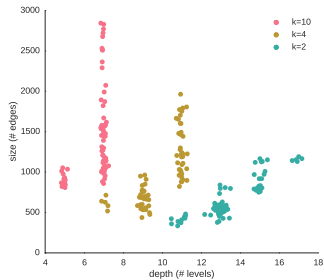
- ⊕ not committing to complex structures too early
- ⊕ same expressive power: successive splits allow for more node children
- ⊕ reducing node out fan increases the depth
- ⊕ same accuracy, smaller networks

Experimental setting

Classical setting for **generative** graphical models structure learning [Gens and Domingos 2013]:

- ⊕ comparing the **average log-likelihood** on predicting instances from a test set
- ⊕ 19 binary datasets from classification, recommendation, frequent pattern mining...[Lowd and Davis 2010] [Haaren and Davis 2012]
- ⊕ Training 75% Validation 10% Test 15% splits (no cv)
- ⊕ Model selection via *grid search* in the same parameter space:
 - ⊗ $\lambda \in \{0.2, 0.4, 0.6, 0.8\}$,
 - ⊗ $\rho \in \{5, 10, 15, 20\}$,
 - ⊗ $m \in \{1, 50, 100, 500\}$,
 - ⊗ $\alpha \in \{0.1, 0.2, 0.5, 1.0, 2.0\}$
- ⊕ comparing our variants against LearnSPN, ID-SPN and MT [Meilă and Jordan 2000]

Depth VS Size




Regularizing by effective early stopping

LearnSPN regularization is governed by α and m , however can be very ineffective:

- ⊕ naive factorizations are too strong assumptions
- ⊕ best likelihood structures prefer smaller values for m to get accurate naive factorizations


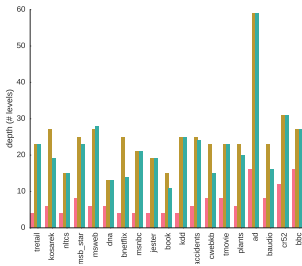
Substituting naive factorizations with Bayesian trees as leaf distributions

$$P(\mathbf{X}) = \prod_j P(X_j | Pa_j):$$



- ⊕ learnable with Chow-Liu algorithm
- ⊕ still tractable multivariate distributions for marginals, conditionals and MPE
- ⊕ same or higher accuracy
- ⊕ less complex structures for larger values of m

Early stopping exp



Strengthening by model averaging

Interpreting sum nodes as *general additive estimators*. Leveraging classic statistical tools to learn them: **bagging**.

Draw k bootstrapped samples from the data, then grow an SPN S_{B_i} on each of them. Join them into a single SPN \hat{S} with a sum node:

$$\hat{S} = \sum_{i=1}^k \frac{1}{k} S_{B_i}$$

More robustness and less variance in the model.

Exponential number of nodes if done for each sum node (bootstrapping only at the root).

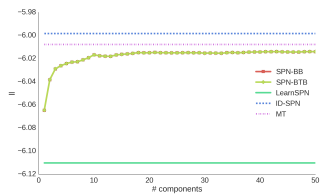
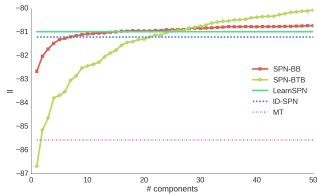
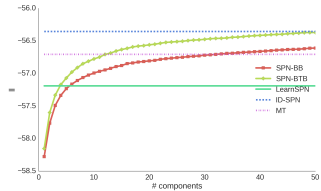
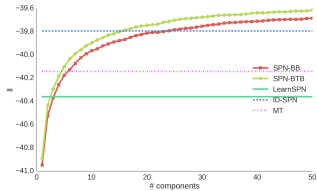
Two variants in the experiments: SPN-BB and SPN-BTB, whether Chow-Liu trees are employed or not.

LL exp

	LearnSPN	SPN-B	SPN-BT	ID-SPN	SPN-BB	SPN-BTB	MT
NLTCS	-6.110	-6.048	-6.048	-5.998	-6.014	-6.014	-6.008
MSNBC	-6.099	-6.040	-6.039	-6.040	-6.032	-6.033	-6.076
KDDCup2k	-2.185	-2.141	-2.141	-2.134	-2.122	-2.121	-2.135
Plants	-12.878	-12.813	-12.683	-12.537	-12.167	-12.089	-12.926
Audio	-40.360	-40.571	-40.484	-39.794	-39.685	-39.616	-40.142
Jester	-53.300	-53.537	-53.546	-52.858	-52.873	-53.600	-53.057
Netflix	-57.191	-57.730	-57.450	-56.355	-56.610	-56.371	-56.706
Accidents	-30.490	-29.342	-29.265	-26.982	-28.510	-28.351	-29.692
Retail	-11.029	-10.944	10.942	-10.846	-10.858	-10.858	-10.836
Pumsb-star	-24.743	-23.315	-23.077	-22.405	-22.866	-22.664	-23.702
DNA	-80.982	-81.913	-81.840	-81.211	-80.730	-80.068	-85.568
Kosarek	-10.894	-10.719	-10.685	-10.599	-10.690	-10.578	-10.615
MSWeb	-10.108	-9.833	-9.838	-9.726	-9.630	-9.614	-9.819
Book	-34.969	-34.306	-34.280	-34.136	-34.366	-33.818	-34.694
EachMovie	-52.615	-51.368	-51.388	-51.512	-50.263	-50.414	-54.513
WebKB	-158.164	-154.283	-153.911	-151.838	-151.341	-149.851	-157.001
Reuters-52	-85.414	-83.349	-83.361	-83.346	-81.544	-81.587	-86.531
BBC	-249.466	-247.301	-247.254	-248.929	-226.359	-226.560	-259.962
Ad	-19.760	-16.234	-15.885	-19.053	-13.785	-13.595	-16.012

Table : Average test log likelihoods for all algorithms.

Bagging exp



Conclusions and Further work

- ⊕ Structure quality evaluation matters
- ⊕ Deeper networks by applying a simplicity bias when splitting
- ⊕ Regularized SPNs by introducing Chow-Liu trees as leaves
- ⊕ More robust and accurate SPNs with bootstrapped sum nodes

References

- ⊕ Dennis, Aaron and Dan Ventura (2012). ``Learning the Architecture of Sum-Product Networks Using Clustering on Variables". In: *Advances in Neural Information Processing Systems 25*. Curran Associates, Inc., pp. 2033–2041 (cit. on p. 7).
- ⊕ Gens, Robert and Pedro Domingos (2013). ``Learning the Structure of Sum-Product Networks". In: *Proceedings of the 30th International Conference on Machine Learning*. JMLR Workshop and Conference Proceedings, pp. 873–880 (cit. on pp. 7, 12).
- ⊕ Haaren, Jan Van and Jesse Davis (2012). ``Markov Network Structure Learning: A Randomized Feature Generation Approach". In: *Proceedings of the 26th Conference on Artificial Intelligence*. AAAI Press (cit. on p. 12).
- ⊕ Lowd, Daniel and Jesse Davis (2010). ``Learning Markov Network Structure with Decision Trees". In: *Proceedings of the 10th IEEE International Conference on Data Mining*. IEEE Computer Society Press, pp. 334–343 (cit. on p. 12).
- ⊕ Martens, James and Venkatesh Medabalimi (2014). ``On the Expressive Efficiency of Sum Product Networks". In: *CoRR* abs/1411.7717 (cit. on p. 8).
- ⊕ Meilä, Marina and Michael I. Jordan (2000). ``Learning with mixtures of trees". In: *Journal of Machine Learning Research* 1, pp. 1–48 (cit. on p. 12).
- ⊕ Peharz, Robert, Bernhard Geiger, and Franz Pernkopf (2013). ``Greedy Part-Wise Learning of Sum-Product Networks". In: *Machine Learning and Knowledge Discovery in Databases*. Vol. 8189. LNCS. Springer, pp. 612–627 (cit. on p. 7).
- ⊕ Rooshenas, Amirmohammad and Daniel Lowd (2014). ``Learning Sum-Product Networks with Direct and Indirect Variable Interactions". In: *Proceedings of the 31st International Conference on Machine Learning*. JMLR Workshop and Conference Proceedings, pp. 710–718 (cit. on p. 7).
- ⊕ Zhao, Han, Mazen Melibari, and Pascal Poupart (2015). ``On the Relationship between Sum-Product Networks and Bayesian Networks". In: *CoRR* abs/1501.01239. URL: <http://arxiv.org/abs/1501.01239> (cit. on p. 8).